Pixel FFN: fix the cost of short CPU bursts

Server-like cadence, real Qwen layers 18-23, packed CPU weights, ADB TCP

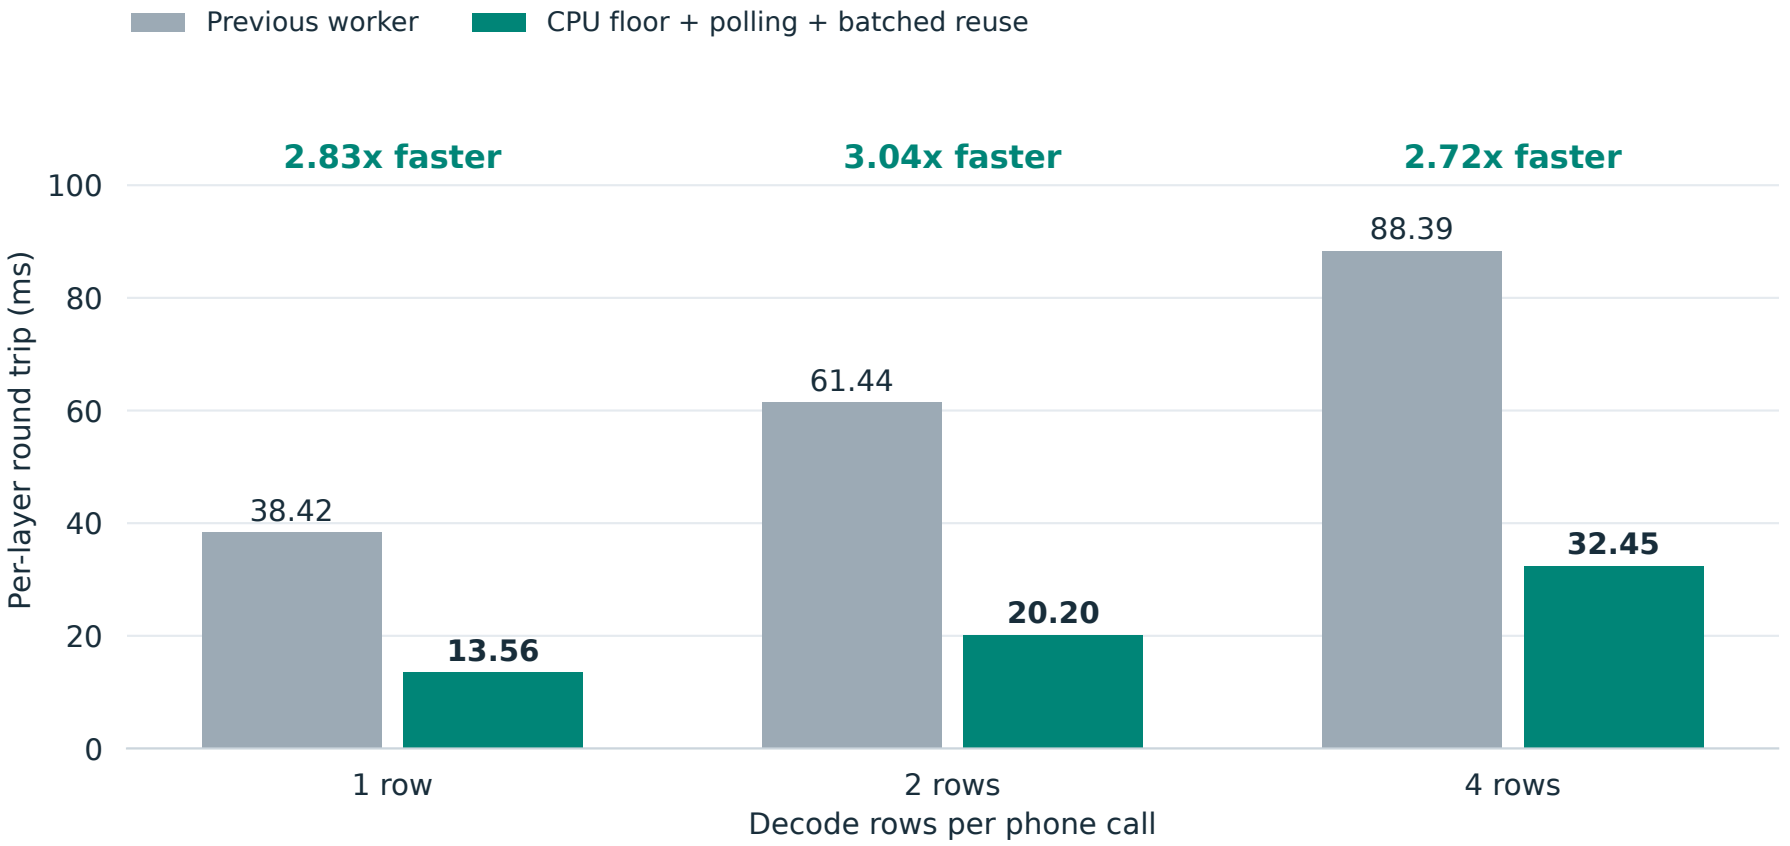

Before: mean of two control medians. After: optimized median; 72 calls per row count.
Byte-identical kernel outputs. These are layer-call latencies, not whole-model token latencies.